

Jamboree Education - Linear Regression

Context

Jamboree has helped thousands of students like you make it to top colleges abroad. Be it GMAT, GRE or SAT, their unique problem-solving methods ensure maximum scores with minimum effort. They recently launched a feature where students/learners can come to their website and check their probability of getting into the IVY league college. This feature estimates the chances of graduate admission from an Indian perspective.

Objective / Problem statement

Your analysis will help Jamboree in understanding what factors are important in graduate admissions and how these factors are interrelated among themselves. It will also help predict one's chances of admission given the rest of the variables.

In [258]:

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6 # Linear regression Library
7 from sklearn.linear_model import LinearRegression, Ridge, Lasso
8 from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
9 from sklearn.model_selection import train_test_split, GridSearchCV, KFold
10 from sklearn.preprocessing import StandardScaler, MinMaxScaler, PolynomialFeatures
11 from sklearn.pipeline import make_pipeline
12
13 #stats model Library
14 import statsmodels.api as sm
15 from statsmodels.stats.outliers_influence import variance_inflation_factor
16
17 # hypothesis testing Library
18 from scipy.stats import shapiro
19
20 # math Library
21 import math
```

In [5]:

```
1 # Load Jamboree education data
2 df = pd.read_csv(r'https://d2beiqkhq929f0.cloudfront.net/public_assets/assets/000/001/839/original/Jamboree_Admission.csv')
```

Target / dependent feature - 'Chance of Admit'

Independent feature - 'GRE Score', 'TOEFL Score', 'University Rating', 'SOP', 'LOR ', 'CGPA', 'Research'

In [6]:

```
1 df.head()
```

Out[6]:

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
0	1	337	118	4	4.5	4.5	9.65	1	0.92
1	2	324	107	4	4.0	4.5	8.87	1	0.76
2	3	316	104	3	3.0	3.5	8.00	1	0.72
3	4	322	110	3	3.5	2.5	8.67	1	0.80
4	5	314	103	2	2.0	3.0	8.21	0	0.65

In [7]:

```
1 df.columns
```

Out[7]:

Index(['Serial No.', 'GRE Score', 'TOEFL Score', 'University Rating', 'SOP', 'LOR ', 'CGPA', 'Research', 'Chance of Admit '], dtype='object')

Null values are not present

In [8]:

```
1 df.isna().sum()
```

Out[8]:

Serial No.0

GRE Score0

TOEFL Score0

University Rating0

SOP0

LOR0

CGPA0

Research0

Chance of Admit0

dtype: int64

- GRE score as highest mean value
- Mean of Chance of Admit is 0.72

In [9]:

```
1 df.describe()
```

Out[9]:

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
count	500.000000	500.000000	500.000000	500.000000	500.000000	500.00000	500.000000	500.000000	500.00000
mean	250.500000	316.472000	107.192000	3.114000	3.374000	3.48400	8.576440	0.560000	0.72174
std	144.481833	11.295148	6.081868	1.143512	0.991004	0.92545	0.604813	0.496884	0.14114
min	1.000000	290.000000	92.000000	1.000000	1.000000	1.00000	6.800000	0.000000	0.34000
25%	125.750000	308.000000	103.000000	2.000000	2.500000	3.00000	8.127500	0.000000	0.63000
50%	250.500000	317.000000	107.000000	3.000000	3.500000	3.50000	8.560000	1.000000	0.72000
75%	375.250000	325.000000	112.000000	4.000000	4.000000	4.00000	9.040000	1.000000	0.82000
max	500.000000	340.000000	120.000000	5.000000	5.000000	5.00000	9.920000	1.000000	0.97000

All features data type are numeric

```
In [10]: 1 df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Serial No.            500 non-null   int64
1   GRE Score              500 non-null   int64
2   TOEFL Score            500 non-null   int64
3   University Rating      500 non-null   int64
4   SOP                    500 non-null   float64
5   LOR                    500 non-null   float64
6   CGPA                   500 non-null   float64
7   Research               500 non-null   int64
8   Chance of Admit        500 non-null   float64
dtypes: float64(4), int64(5)
memory usage: 35.3 KB
```

Univariant analysis

```
In [31]: 1 col = ['GRE Score', 'TOEFL Score', 'University Rating', 'SOP', 'LOR ', 'CGPA', 'Research', 'Chance of Admit ']
2 for i in col:
3     print('Unique values of',i, 'column =', df[i].nunique())
4     print('-----')

Unique values of GRE Score column = 49
-----
Unique values of TOEFL Score column = 29
-----
Unique values of University Rating column = 5
-----
Unique values of SOP column = 9
-----
Unique values of LOR  column = 9
-----
Unique values of CGPA column = 184
-----
Unique values of Research column = 2
-----
Unique values of Chance of Admit  column = 61
-----
```

Unique value counts

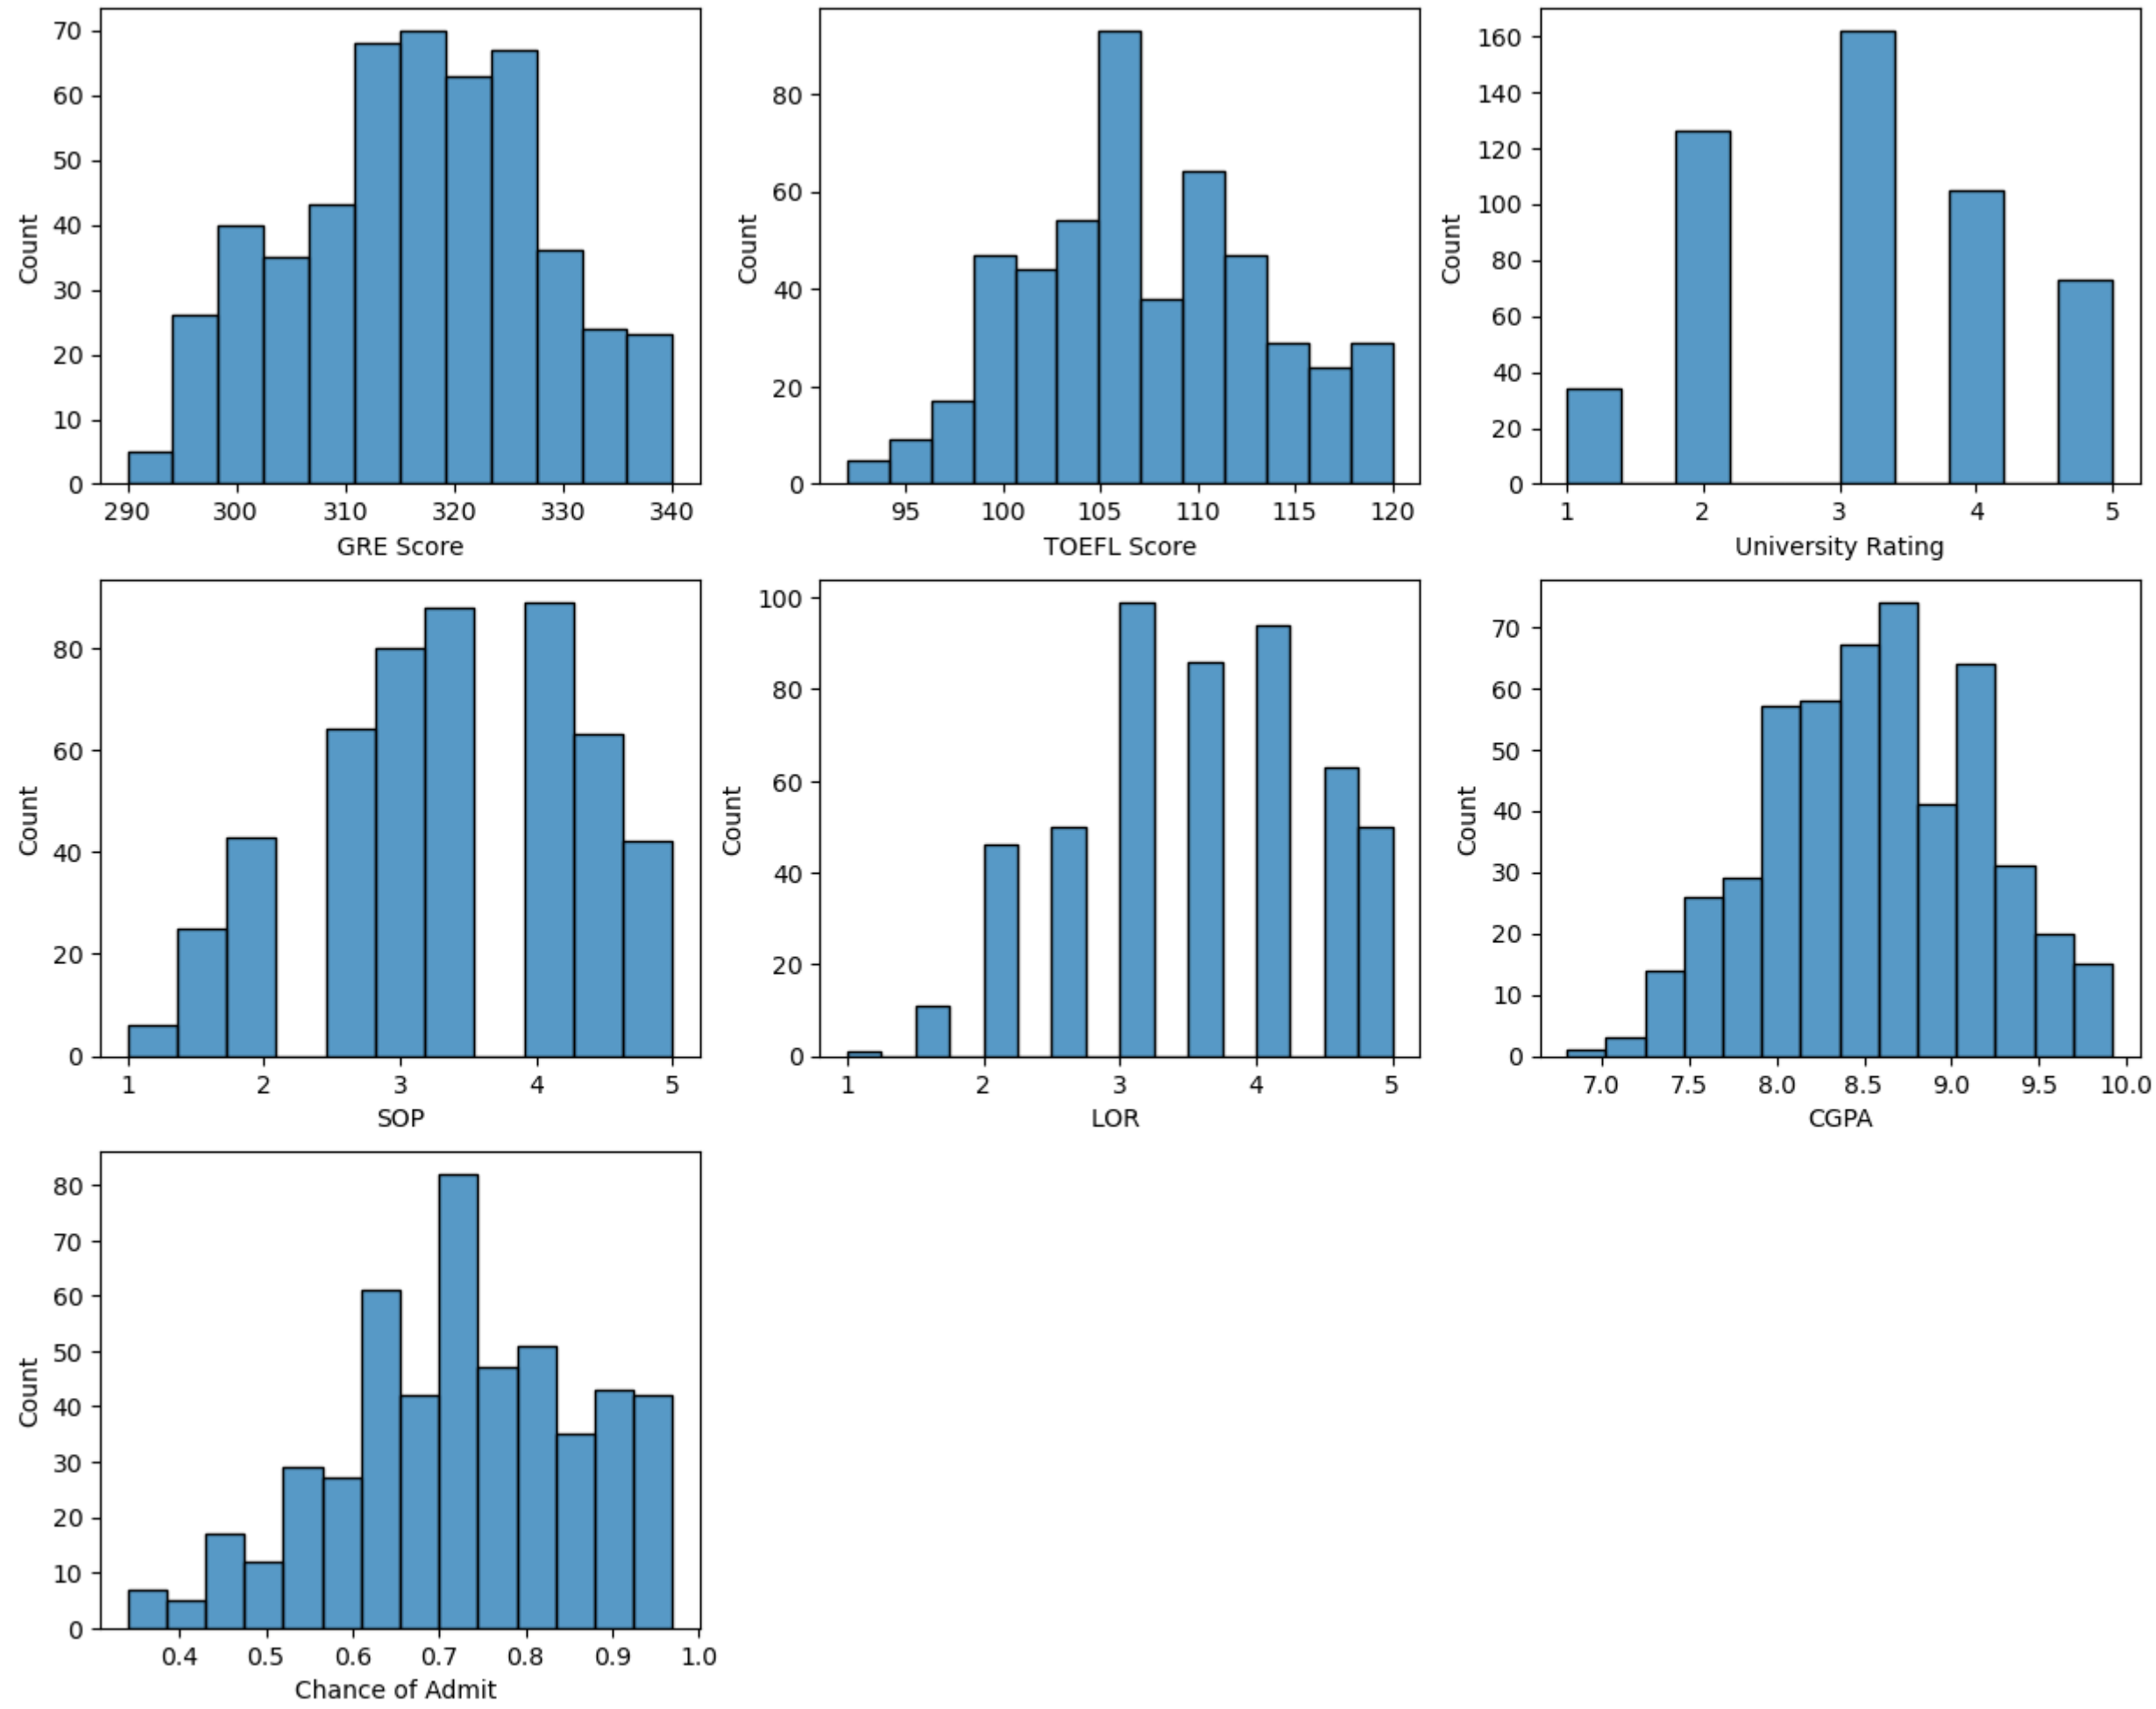
- University Rating --> 3 as highest rating
- SOP --> 4 as highest rating
- LOR --> 3 as highest rating
- Research --> 1 is majority/

```
In [140]: 1 col = ['University Rating', 'SOP', 'LOR ', 'Research']
2 for i in col:
3     print('Unique values of',i, 'column')
4     print((df[i].value_counts()/df.shape[0]) * 100)
5     print('-----')

Unique values of University Rating column
3      32.4
2      25.2
4      21.0
5      14.6
1       6.8
Name: University Rating, dtype: float64
-----
Unique values of SOP column
4.0      17.8
3.5      17.6
3.0      16.0
2.5      12.8
4.5      12.6
2.0       8.6
5.0       8.4
1.5       5.0
1.0       1.2
Name: SOP, dtype: float64
-----
Unique values of LOR  column
3.0      19.8
4.0      18.8
3.5      17.2
4.5      12.6
2.5      10.0
5.0      10.0
2.0       9.2
1.5       2.2
1.0       0.2
Name: LOR , dtype: float64
-----
Unique values of Research column
1      56.0
0      44.0
Name: Research, dtype: float64
-----
```

Histogram

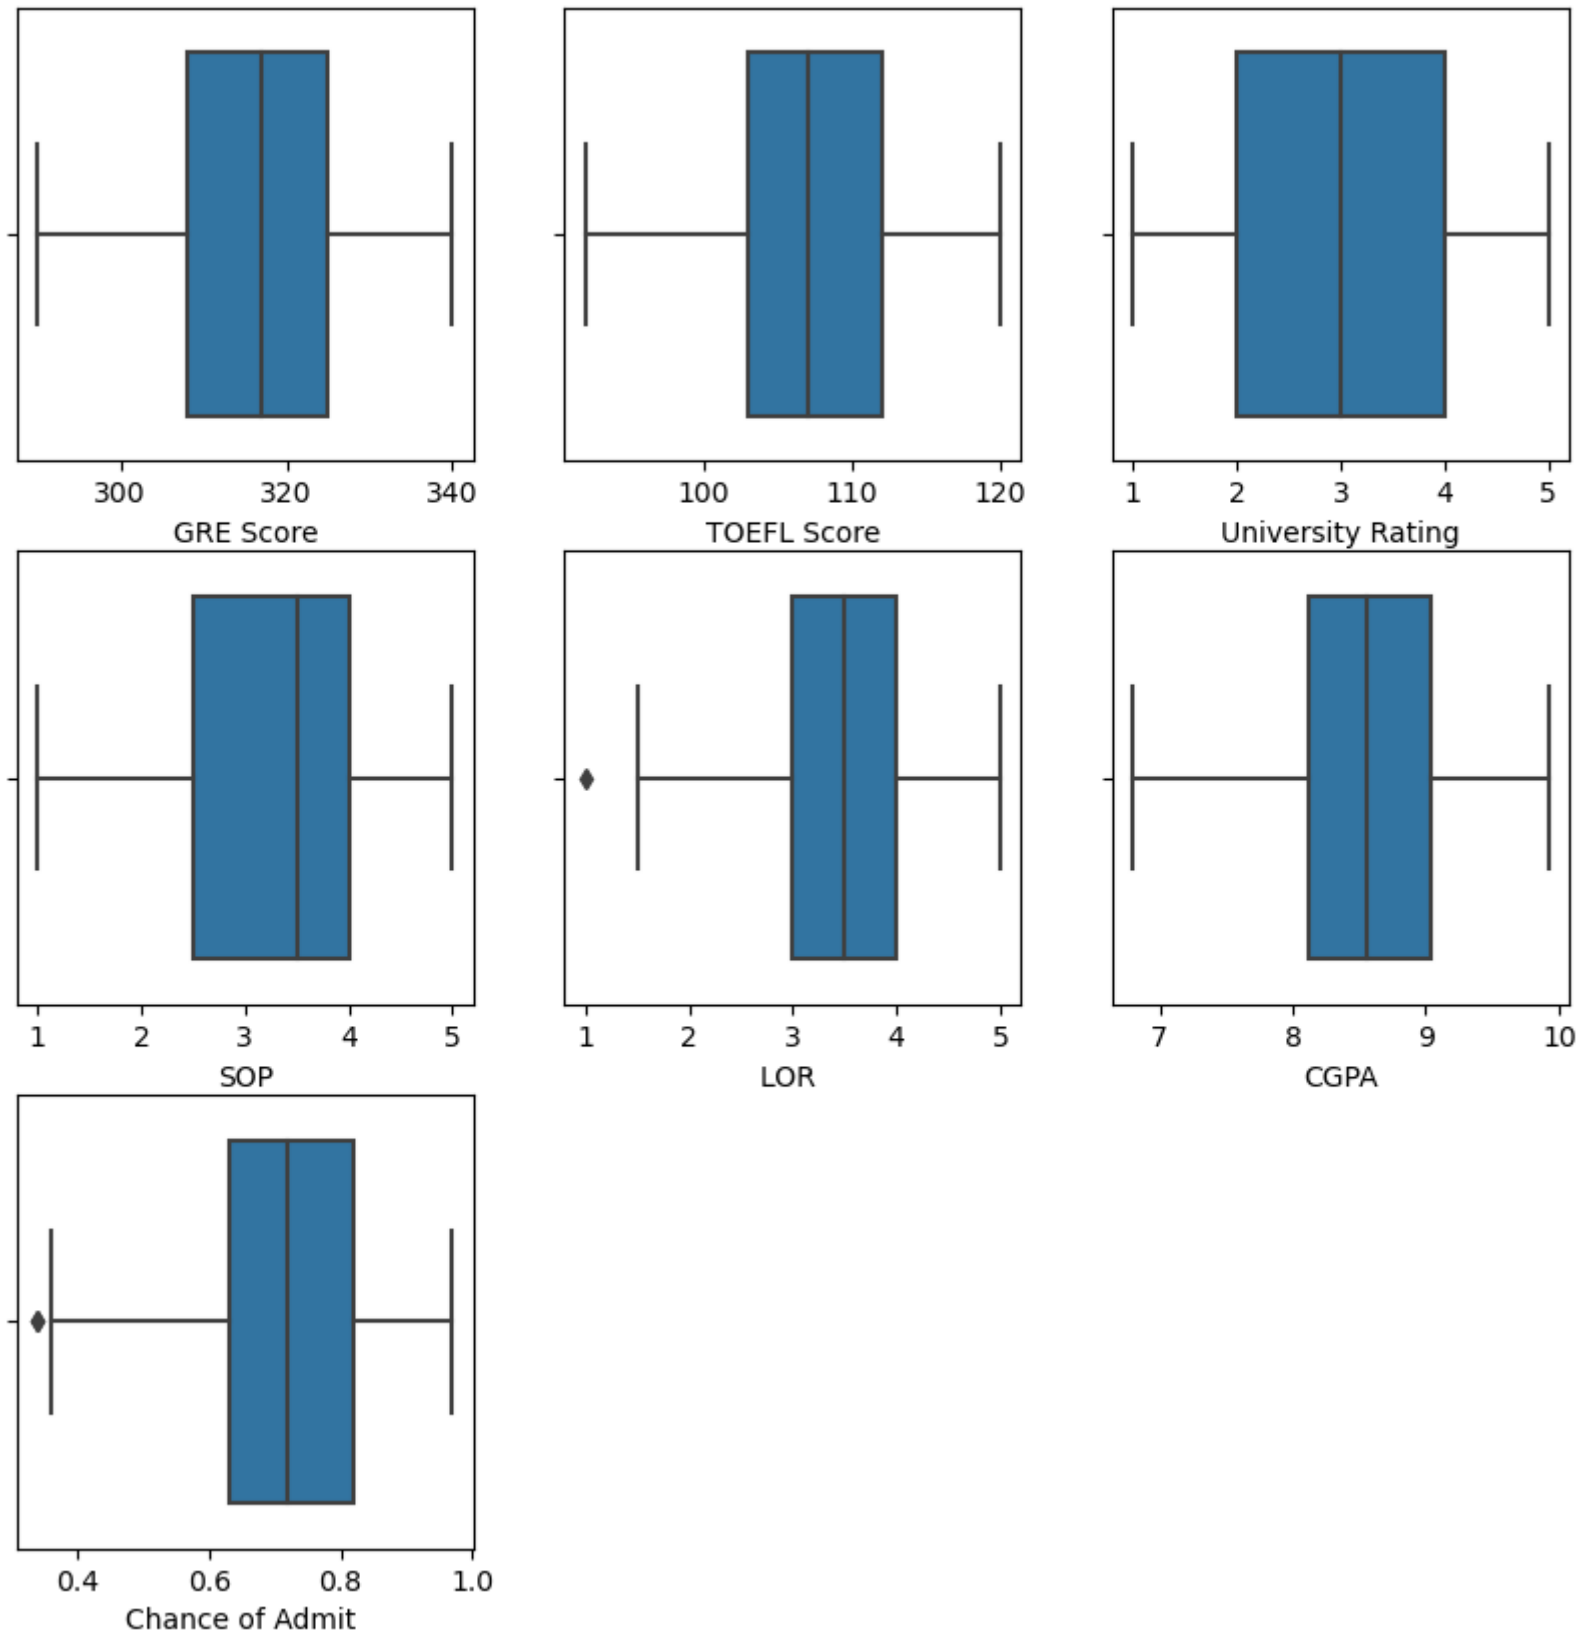
```
In [101]: 1 col = df.columns.drop (['Serial No.','Research'])
2 fig, ax = plt.subplots(3,3, figsize = (15,12))
3 ax[2,1].set_axis_off()
4 ax[2,2].set_axis_off()
5 c = 0
6 for i in range(3):
7     for j in range(3):
8         if i == 2 and j==1:
9             break
10        sns.histplot(ax = ax[i,j], x = df[col[c]])
11        c = c+1
```



Outlier detection

- LOR and Chance of admit features as few outliers and remaining features doesnt have outlier

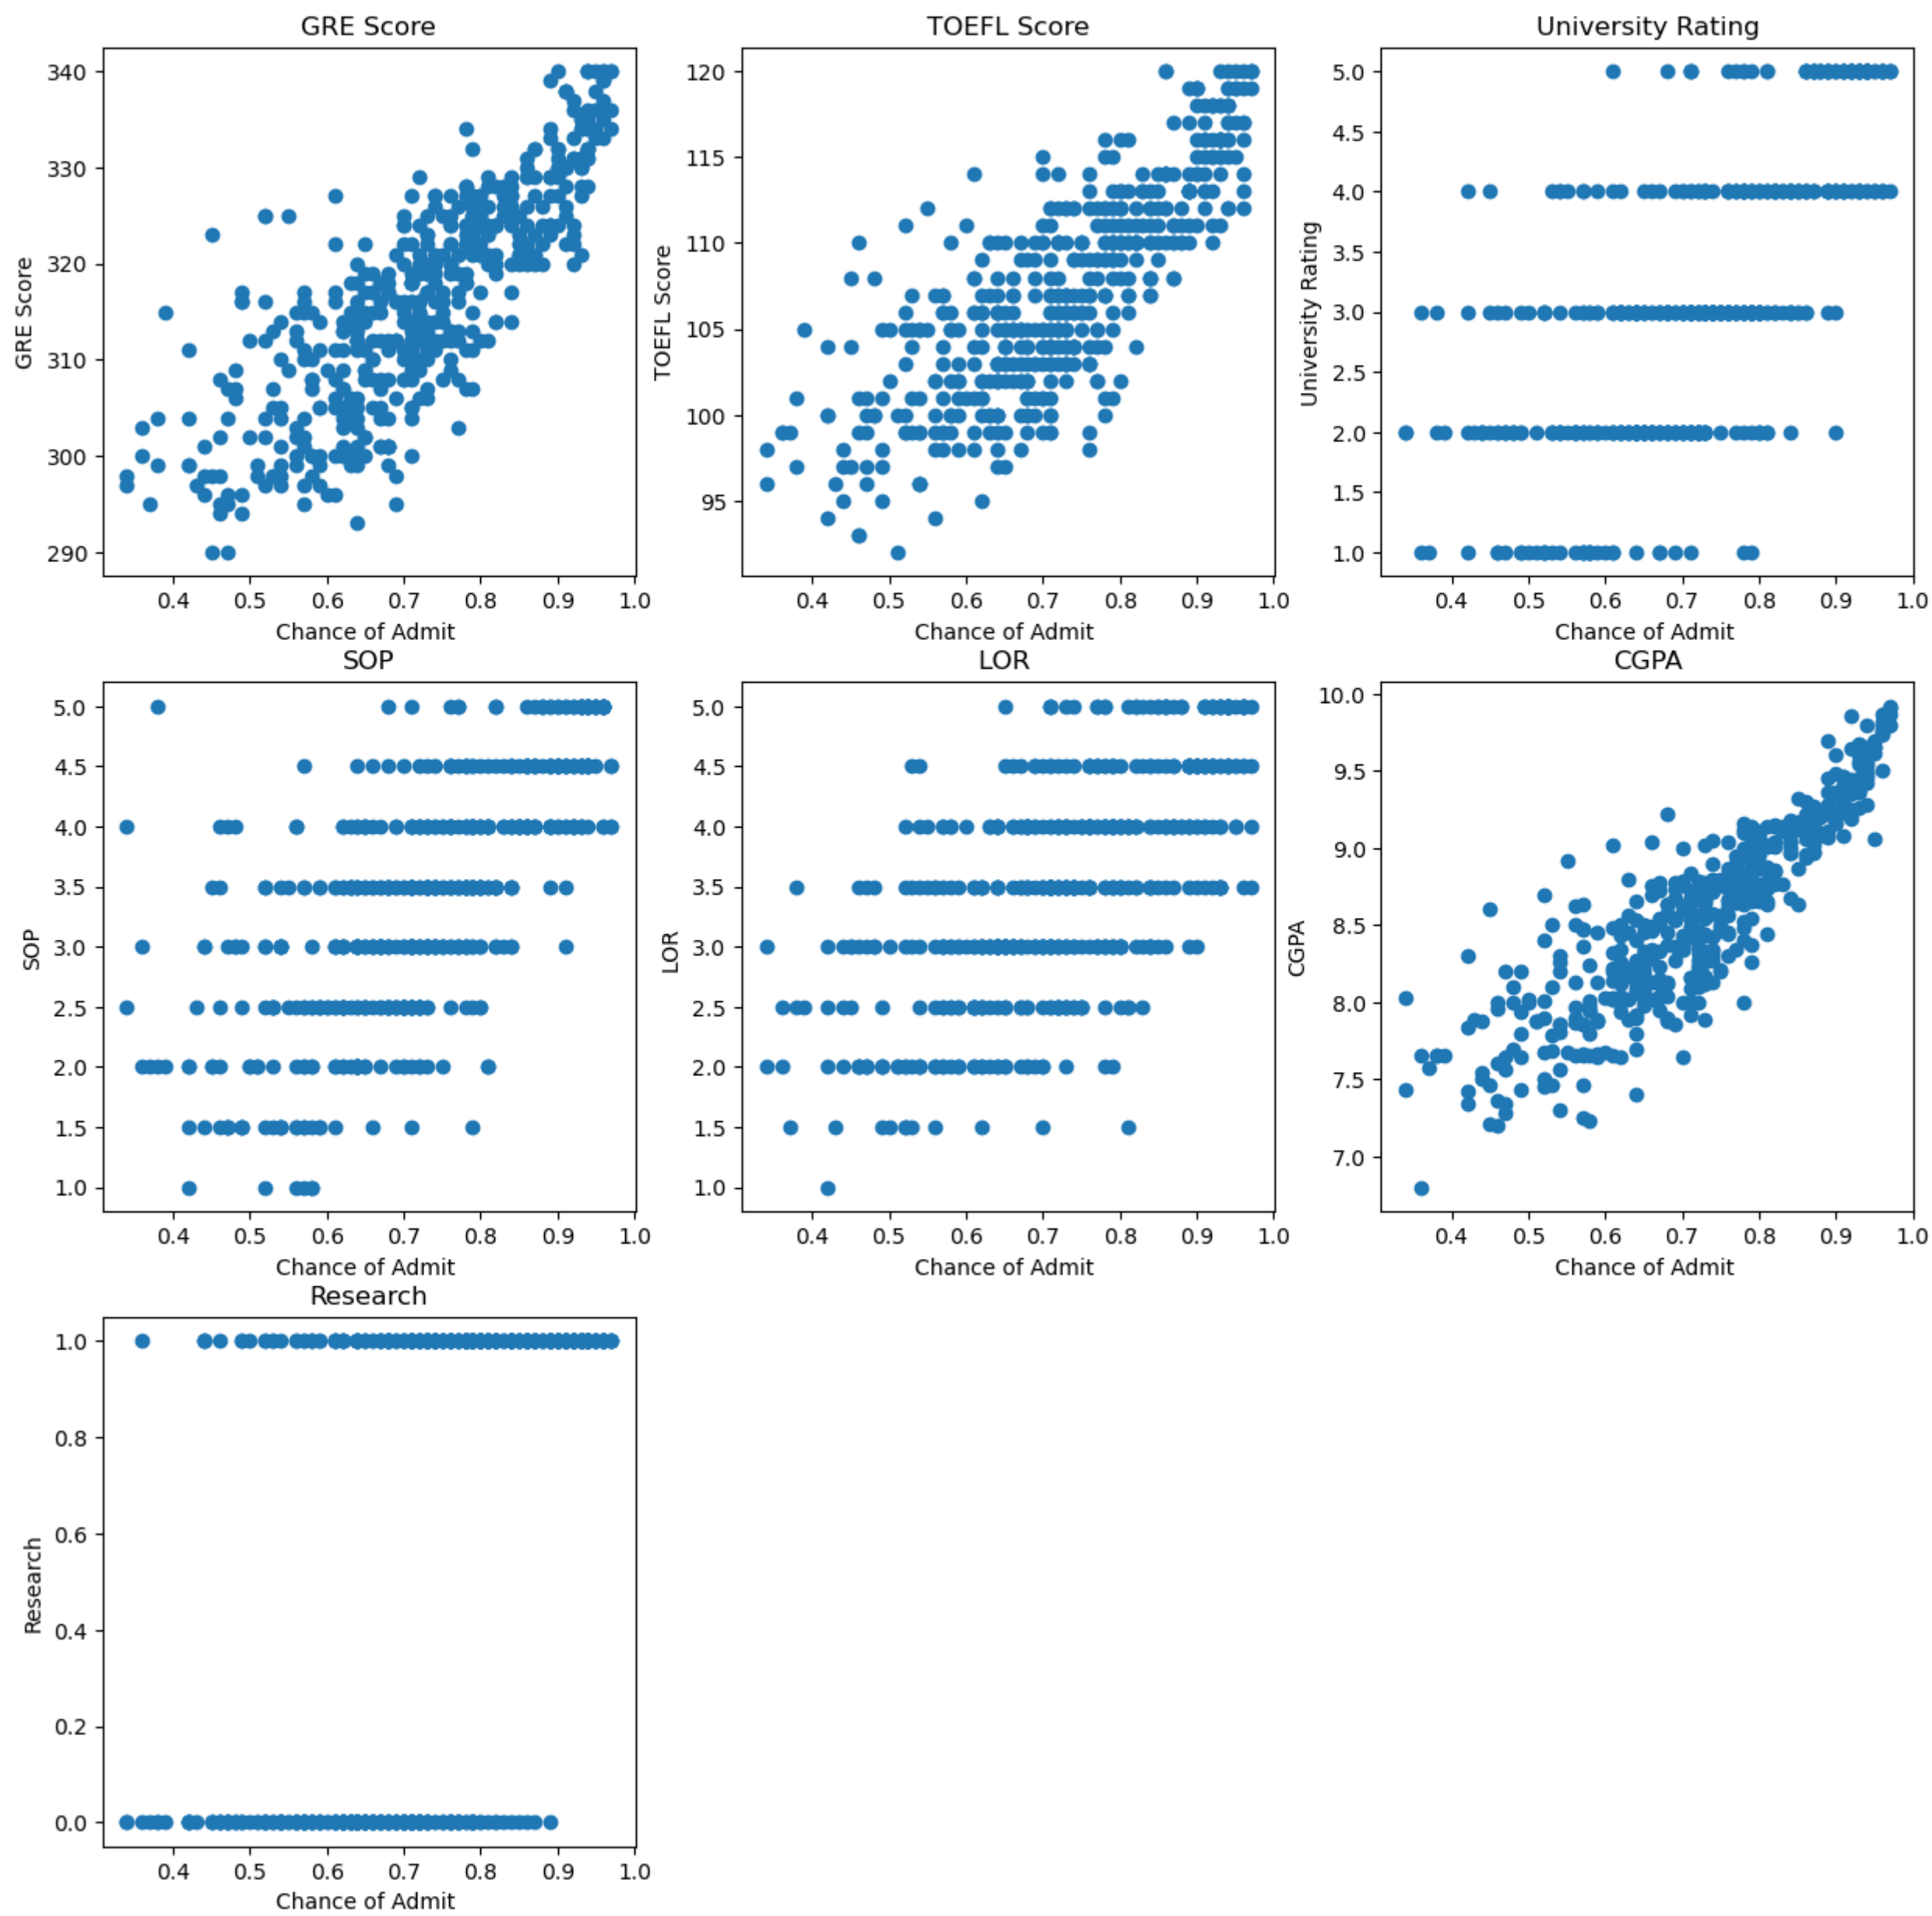
```
In [88]: 1 col = df.columns.drop (['Serial No.','Research'])
2 fig, ax = plt.subplots(3,3, figsize = (10,10))
3 ax[2,1].set_axis_off()
4 ax[2,2].set_axis_off()
5 c = 0
6 for i in range(3):
7     for j in range(3):
8         if i == 2 and j==1:
9             break
10        sns.boxplot(ax = ax[i,j], x = df[col[c]])
11        c = c+1
12
```



Bivariate Analysis

```
In [134]: 1 col = df.columns.drop (['Serial No.', 'Chance of Admit '])
2 fig, ax = plt.subplots(3,3, figsize = (15,15))
3 fig.suptitle('Bivariate Analysis')
4 ax[2,1].set_axis_off()
5 ax[2,2].set_axis_off()
6 c = 0
7
8 for i in range(3):
9     for j in range(3):
10         if i == 2 and j==1:
11             break
12         ax[i,j].scatter(x = df['Chance of Admit '], y = df[col[c]])
13         ax[i,j].set_xlabel("Chance of Admit ")
14         ax[i,j].set_ylabel(col[c])
15         ax[i,j].set_title(col[c])
16         c = c +1
```

Bivariate Analysis




Bivariate insights :

- Independent features are linear correlated to Target features as per above scatter plot

```
In [135]: 1 sns.heatmap(df.corr(method='pearson'), annot=True)
```

Out[135]: <Axes: >



Heatmap insights :

- Independent features are linear correlated to Target features

Linear regression Assumption checking

```
In [184]: 1 data = df.drop(columns = ['Serial No.'])
2 col_names = data.columns
3
4 # scale the model for checking assumption
5 scale = StandardScaler()
6 data_scale = scale.fit_transform(data)
7 data_scale = pd.DataFrame(data_scale, columns = col_names)
8 X_asm_sc = data_scale.drop(columns = ['Chance of Admit '])
9 y_asm_sc = data_scale['Chance of Admit ']
```

Multicollinearity check by VIF score

- As per Vif values of each column there are not multicollinearity between independent columns as All VIF values as below 5.

```
In [185]: 1 # create dataframe to VIF values
2 vif = pd.DataFrame()
3
4 # create features column for comparision
5 vif['Features'] = X_asm_sc.columns
6
7 # create VIF values for all independent columns - GRE Score, TOEFL Score, University Rating, SOP, LOR, CGPA, 'Research'
8 vif['Vif values'] = [variance_inflation_factor(X_asm_sc.values, i) for i in range(X_asm_sc.shape[1])]
9
10 # round values
11 vif['Vif values'] = round(vif['Vif values'],2)
12
13 # sort by values in decending order
14 vif= vif.sort_values(by = 'Vif values', ascending = False )
15 vif
```

Out[185]:

	Features	Vif values
5	CGPA	4.78
0	GRE Score	4.46
1	TOEFL Score	3.90
3	SOP	2.84
2	University Rating	2.62
4	LOR	2.03
6	Research	1.49

Error should be normal distributed

- As per checking with graphical and non graphical given data's error is not normally distributed
- Error is left skewed.


```
In [190]: 1 # adding intercept values to data frame when we run OLS model from Statsmodels Library
2 X_ecc = sm.add_constant(X_asn_sc)
3
4 # Train the model
5 sm_model = sm.OLS(y_asn_sc, X_ecc).fit()
6
7 # Get results
8 print(sm_model.summary())
9
10 # get y_prediction values
11 y_pre = sm_model.predict(X_ecc)
12
13 # Error = y - y_pre
14 errors = sm_model.resid
15
16 # plot histogram
17 sns.histplot(errors)
```


OLS Regression Results						
=====						
Dep. Variable:	Chance of Admit	R-squared:	0.822			
Model:	OLS	Adj. R-squared:	0.819			
Method:	Least Squares	F-statistic:	324.4			
Date:	Mon, 25 Dec 2023	Prob (F-statistic):	8.21e-180			
Time:	19:55:23	Log-Likelihood:	-278.12			
No. Observations:	500	AIC:	572.2			
Df Residuals:	492	BIC:	605.9			
Df Model:	7					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	-3.322e-16	0.019	-1.75e-14	1.000	-0.037	0.037
GRE Score	0.1487	0.040	3.700	0.000	0.070	0.228
TOEFL Score	0.1197	0.038	3.184	0.002	0.046	0.194
University Rating	0.0481	0.031	1.563	0.119	-0.012	0.109
SOP	0.0111	0.032	0.348	0.728	-0.052	0.074
LOR	0.1105	0.027	4.074	0.000	0.057	0.164
CGPA	0.5073	0.042	12.198	0.000	0.426	0.589
Research	0.0856	0.023	3.680	0.000	0.040	0.131
=====						
Omnibus:	112.770	Durbin-Watson:	0.796			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	262.104			
Skew:	-1.160	Prob(JB):	1.22e-57			
Kurtosis:	5.684	Cond. No.	5.65			
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
Out[190]: <Axes: ylabel='Count'>
```




```
In [193]: 1 # Shapiro test to check normality of errors
2 # null hypothesis - data is normally distributed
3 # Alternate hypothesis - data is not normally distributed
4 # signigicance level, alpha = 0.05
5 alp = 0.05 # alpha value
6 statistic , p_value = shapiro(errors)
7
8 if p_value > alp:
9     print(f'Pvalue : {p_value}')
10    print('Errors is normally distributed')
11 else:
12     print(f'Pvalue : {p_value}')
13    print('Errors is not normally distributed')
```

Pvalue : 4.826223479832605e-15
Errors is not normally distributed

Checking for Heteroscedasticity

- As per below plot, errors looks like having heteriscedasticity
- As per mean of residual = values is almost zero but required '0'

```
In [199]: 1 heter = sns.scatterplot(x = y_pre, y = errors)
2 heter.axhline(0, color = 'red', linewidth = 2)
3 plt.show()
```



```
In [200]: 1 # mean of residual
2 errors.mean()
```

Out[200]: -1.971756091734278e-15

Split data - Train and Test

```
In [201]: 1 X_train, X_test, y_train, y_test = train_test_split(X_asm_sc, y_asm_sc, test_size = 0.2, random_state = 1)
2 X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

Out[201]: ((400, 7), (100, 7), (400,), (100,))

Linear regression by stats model

- Train data:
 - R2 and Adj R2 score are 0.822 and 0.818 respectively
 - SOP and University Rating columns probability is more than 0.05. Where feature is not important as Null hypothesis testing.
 - GRE Score, TOEFL Score, LOR, CGPA and Research features are important as per probability values.
 - 95 percentile values lie in the range for each column mentioned.

- Test data:
 - R2 and Adj R2 score are 0.82 and 0.81 respectively
 - MAE score test : 0.29
 - RMSE score test : 0.42

- Model prediction is almost same in Train and Test data. so model is best fit

- Adjusted R2 score is about 0.81
- R2 score is about 0.8 ~ 0.82
- Model is predicts 80% confident

```
In [203]: 1 # by Stats model- Linear regression
2 sm_X_train = sm.add_constant(X_train)
3 sm_model_train = sm.OLS(y_train,sm_X_train).fit()
4
5 # summary
6 print(sm_model_train.summary())
```

OLS Regression Results						
=====						
Dep. Variable:	Chance of Admit	R-squared:	0.822			
Model:	OLS	Adj. R-squared:	0.818			
Method:	Least Squares	F-statistic:	257.7			
Date:	Mon, 25 Dec 2023	Prob (F-statistic):	2.10e-142			
Time:	20:32:32	Log-Likelihood:	-224.33			
No. Observations:	400	AIC:	464.7			
Df Residuals:	392	BIC:	496.6			
Df Model:	7					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	0.0081	0.021	0.377	0.707	-0.034	0.050
GRE Score	0.1466	0.047	3.135	0.002	0.055	0.239
TOEFL Score	0.1368	0.043	3.156	0.002	0.052	0.222
University Rating	0.0497	0.036	1.387	0.166	-0.021	0.120
SOP	0.0211	0.036	0.591	0.555	-0.049	0.091
LOR	0.0946	0.030	3.105	0.002	0.035	0.154
CGPA	0.5001	0.047	10.743	0.000	0.409	0.592
Research	0.0700	0.026	2.668	0.008	0.018	0.122
=====						
Omnibus:	80.594	Durbin-Watson:	1.932			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	167.116			
Skew:	-1.064	Prob(JB):	5.14e-37			
Kurtosis:	5.346	Cond. No.	5.92			
=====						

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [268]:

```
1 # with test data checking
2 sm_X_test = sm.add_constant(X_test)
3 sm_y_pre = sm_model_train.predict(sm_X_test)
4
5 # R2 score
6 r2_test = r2_score(y_test, sm_y_pre)
7 print(f'R2 score test : {round(r2_test,2)}')
8
9 # adjusted R2 score
10 adj_r2 = 1 - (1-r2_test)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
11 print(f'Adjusted R2 score test : {round(adj_r2,2)}')
12
13 # MAE
14 mae_sm = mean_absolute_error(y_test, sm_y_pre)
15 print(f'MAE score test : {round(mae_sm,2)}')
16
17 # RMSE
18 rmse_sm = mean_squared_error(y_test, sm_y_pre)
19 print(f'RMSE score test : {round(math.sqrt(rmse_sm),2)}')
```

R2 score test : 0.82
Adjusted R2 score test : 0.81
MAE score test : 0.29
RMSE score test : 0.42

L1 / Lasso Regularization Regression

L2 / Ridge Regularization Regression

L1/Lasso Regularization regresson

- Train :
 - R2 score = 0.8
 - Adjusted R2 score = 0.78
 - MAE score train : 0.3
 - RMSE score train : 0.42
- Test :
 - R2 score = 0.82
 - Adjusted R2 score = 0.81
 - MAE score test : 0.29
 - RMSE score test : 0.42

- Adjusted R2 score is about 0.81
- R2 score is about 0.8 ~ 0.8
- Model is predicts 80% coinfident

In [266]:

```
1 # L1 / Lasso Regularization Regression
2 def ploy_scale(degree=2, alpha=1.0):
3     return make_pipeline(PolynomialFeatures(degree), Lasso(alpha = alpha))
4
5 param_grid = {'polynomialfeatures__degree': np.arange(1,5), 'lasso__alpha' : [0.01, 0.1, 1 ,1.5, 10]}
6 lasso_r2 = GridSearchCV(estimator = ploy_scale(), param_grid = param_grid, cv = KFold(n_splits = 5), scoring = 'r2')
7
8 # train
9 lasso_r2.fit(X_train, y_train)
10
11 # best degree and alpha
12 print(lasso_r2.best_params_)
13
14 # R2 score
15 print(f'R2 score for train data :{round(lasso_r2.best_score_,2)}')
```

16

```
17 # adjusted R2 score
18 adj_r2_lasso_train = 1 - (1-lasso_r2.best_score_)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
19 print(f'Adjusted R2 score train : {round(adj_r2_lasso_train,2)}')
```

20

```
21 # MAE
22 y_pre_lasso_train = lasso_r2.predict(X_train)
23 mae_lasso = mean_absolute_error(y_train, y_pre_lasso_train)
24 print(f'MAE score train : {round(mae_lasso,2)}')
```

25

```
26 # RMSE
27 rmse_lasso = mean_squared_error(y_train, y_pre_lasso_train)
28 print(f'RMSE score train : {round(math.sqrt(rmse_lasso),2)}')
```

{'lasso__alpha': 0.01, 'polynomialfeatures__degree': 1}
R2 score for train data :0.8
Adjusted R2 score train : 0.78
MAE score train : 0.3
RMSE score train : 0.42

In [265]:

```
1 # Test data
2 y_pre_lasso = lasso_r2.predict(X_test)
3
4 # r2 score
5 r2_test_lasso = r2_score(y_test, y_pre_lasso)
6 print(f'R2 score test : {round(r2_test_lasso,2)}')
```

7

```
8 # adjusted R2 score
9 adj_r2_lasso = 1 - (1-r2_test_lasso)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
10 print(f'Adjusted R2 score test : {round(adj_r2_lasso,2)}')
```

11

```
12 # MAE
13 mae_lasso = mean_absolute_error(y_test, y_pre_lasso)
14 print(f'MAE score test : {round(mae_lasso,2)}')
```

15

```
16 # RMSE
17 rmse_lasso = mean_squared_error(y_test, y_pre_lasso)
18 print(f'RMSE score test : {round(math.sqrt(rmse_lasso),2)}')
```

R2 score test : 0.82
Adjusted R2 score test : 0.81
MAE score test : 0.29
RMSE score test : 0.42

L2 / Ridge Regularization regresson

- Train :
 - R2 score = 0.8
 - Ajusted R2 score = 0.79
 - MAE score train : 0.3
 - RMSE score train : 0.43
- Test :

- R2 score = 0.82
- Ajusted R2 score = 0.8
- MAE score test : 0.29
- RMSE score test : 0.42

- Adjusted R2 score is about 0.8
- R2 score is about 0.8 ~ 0.82
- Model is predicts 80% coinfident

In [267]:

```
1  # L2 / Ridge Regularization Regression
2  def ploy_scale(degree=2, alpha=1.0):
3      return make_pipeline(PolynomialFeatures(degree), Ridge(alpha = alpha))
4
5  param_grid = {'polynomialfeatures__degree': np.arange(1,5), 'ridge__alpha' : [0.01, 0.1, 1 ,1.5, 10, 20]}
6  ridge_r2 = GridSearchCV(estimator = ploy_scale(), param_grid = param_grid, cv = KFold(n_splits = 5), scoring = 'r2')
7
8  # train
9  ridge_r2.fit(X_train, y_train)
10
11 # best degree and alpha
12 print(ridge_r2.best_params_)
13
14 # R2 score
15 print(f'R2 score for train data :{round(ridge_r2.best_score_,2)}')
16
17 # adjusted r2 score
18 adj_r2_ridge_train = 1 - (1-ridge_r2.best_score_)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
19 print(f'Adjusted R2 score test : {round(adj_r2_ridge_train,2)}')
20
21 # MAE
22 y_pre_ridge_train = ridge_r2.predict(X_train)
23 mae_ridge = mean_absolute_error(y_train, y_pre_ridge_train)
24 print(f'MAE score train : {round(mae_ridge,2)}')
25
26 # RMSE
27 rmse_ridge = mean_squared_error(y_train, y_pre_ridge_train)
28 print(f'RMSE score train : {round(math.sqrt(rmse_ridge),2)}')
```

{'polynomialfeatures__degree': 1, 'ridge__alpha': 20}
R2 score for train data :0.8
Adjusted R2 score test : 0.79
MAE score train : 0.3
RMSE score train : 0.43

In [261]:

```
1  # Test data
2  y_pre_ridge = ridge_r2.predict(X_test)
3
4  # r2 score
5  r2_test_ridge = r2_score(y_test, y_pre_ridge)
6  print(f'R2 score test : {round(r2_test_lasso,2)}')
7
8  # adjusted R2 score
9  adj_r2_ridge = 1 - (1-r2_test_ridge)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
10 print(f'Adjusted R2 score test : {round(adj_r2_ridge,2)}')
11
12 # MAE
13 mae_ridge = mean_absolute_error(y_test, y_pre_ridge)
14 print(f'MAE score test : {round(mae_ridge,2)}')
15
16 # RMSE
17 rmse_ridge = mean_squared_error(y_test, y_pre_ridge)
18 print(f'RMSE score test : {round(math.sqrt(rmse_ridge),2)}')
```

R2 score test : 0.82
Adjusted R2 score test : 0.8
MAE score test : 0.29
RMSE score test : 0.42

Insights :

- By all 3 methods, test R2 score -> 0.8 ~ 0.82
- Model is with best fit
- SOP and University Rating columns probability is more than 0.05. Where feature is not important as Null hypothesis testing.
- GRE Score, TOEFL Score, LOR, CGPA and Research features are important as per probability values.
- CGPA is more important when compared to other features.

Recommendations

- for Admition to university - CGPA, GRE and Toefl scores - these are to be seen where student will get selected not.
- Additional training / coaching for these exam CGPA, GRE and Toefl- changes of getting selction will increase.
- Remaining features like SOP, University rating, LOR and research experience will add advantage if available
- If any student is interested in Research, provide intership opportunity with companies.
- Share list of university list for students, so students can prepare for clearing exams.

=====

Dep. Variable:Chance of Admit

R-squared:0.822

Model:OLS

Adj. R-squared:0.818

Method:Least Squares

F-statistic:257.7

Date:Mon, 25 Dec 2023

Prob (F-statistic):2.10e-142

Time:20:32:32

Log-Likelihood:-224.33

No. Observations:400

AIC:464.7

Df Residuals:392

BIC:496.6

Df Model:7

Covariance Type:nonrobust

=====

	coef	std err	t	P> t	[0.025	0.975]
const	0.0081	0.021	0.377	0.707	-0.034	0.050
GRE Score	0.1466	0.047	3.135	0.002	0.055	0.239
TOEFL Score	0.1368	0.043	3.156	0.002	0.052	0.222
University Rating	0.0497	0.036	1.387	0.166	-0.021	0.120
SOP	0.0211	0.036	0.591	0.555	-0.049	0.091
LOR	0.0946	0.030	3.105	0.002	0.035	0.154
CGPA	0.5001	0.047	10.743	0.000	0.409	0.592
Research	0.0700	0.026	2.668	0.008	0.018	0.122

=====

